

# EXPERIMENT 1

**AIM:** Introduction to Google collab and libraries Panda and Numpy

## THEORY:

### 1. Google Colab:

Google Colaboratory, commonly called Google Colab, is a free, cloud-based platform provided by Google that allows users to write and execute Python code in a web browser. It is widely used for machine learning, data analysis, and education because it provides free access to powerful computing resources, including GPUs and TPUs. One of its main advantages is that no local setup is needed — everything runs in the cloud, and notebooks can be easily shared via Google Drive, just like Google Docs.

### 2. Pandas:

Pandas is a popular open-source Python library used for data analysis and manipulation. It provides powerful data structures like:

Series (one-dimensional data)

DataFrame (two-dimensional, tabular data like spreadsheets or SQL tables)

Pandas makes it easy to clean, organize, filter, and analyze datasets, with built-in functions for handling missing values, grouping data, merging datasets, and performing descriptive statistics. It is one of the most essential tools in data science.

### 3. NumPy:

NumPy (Numerical Python) is a fundamental Python library for scientific computing. It provides support for large, multi-dimensional arrays and matrices, along with a wide collection of mathematical functions to operate on them efficiently.

Key features include:

Fast array computations

Linear algebra operations

Fourier transforms

Random number generation

NumPy forms the backbone of many other Python libraries, including Pandas, SciPy, and TensorFlow.

**CODE:**

```
|  
| import pandas as pd  
| import numpy as np  
| import matplotlib.pyplot as plt  
| import seaborn as sns  
| %matplotlib inline  
| from sklearn.model_selection import train_test_split  
  
Double-click (or enter) to edit  
  
▶ diabetes=pd.read_csv('diabetes.csv')  
diabetes.shape  
→ (768, 9)  
  
[ ] diabetes.head()  
→  


|   | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI  | DiabetesPedigreeFunction | Age | Outcome |
|---|-------------|---------|---------------|---------------|---------|------|--------------------------|-----|---------|
| 0 | 6           | 148     | 72            | 35            | 0       | 33.6 | 0.627                    | 50  | 1       |
| 1 | 1           | 85      | 66            | 29            | 0       | 26.6 | 0.351                    | 31  | 0       |
| 2 | 8           | 183     | 64            | 0             | 0       | 23.3 | 0.672                    | 32  | 1       |
| 3 | 1           | 89      | 66            | 23            | 94      | 28.1 | 0.167                    | 21  | 0       |
| 4 | 0           | 137     | 40            | 35            | 168     | 43.1 | 2.288                    | 33  | 1       |

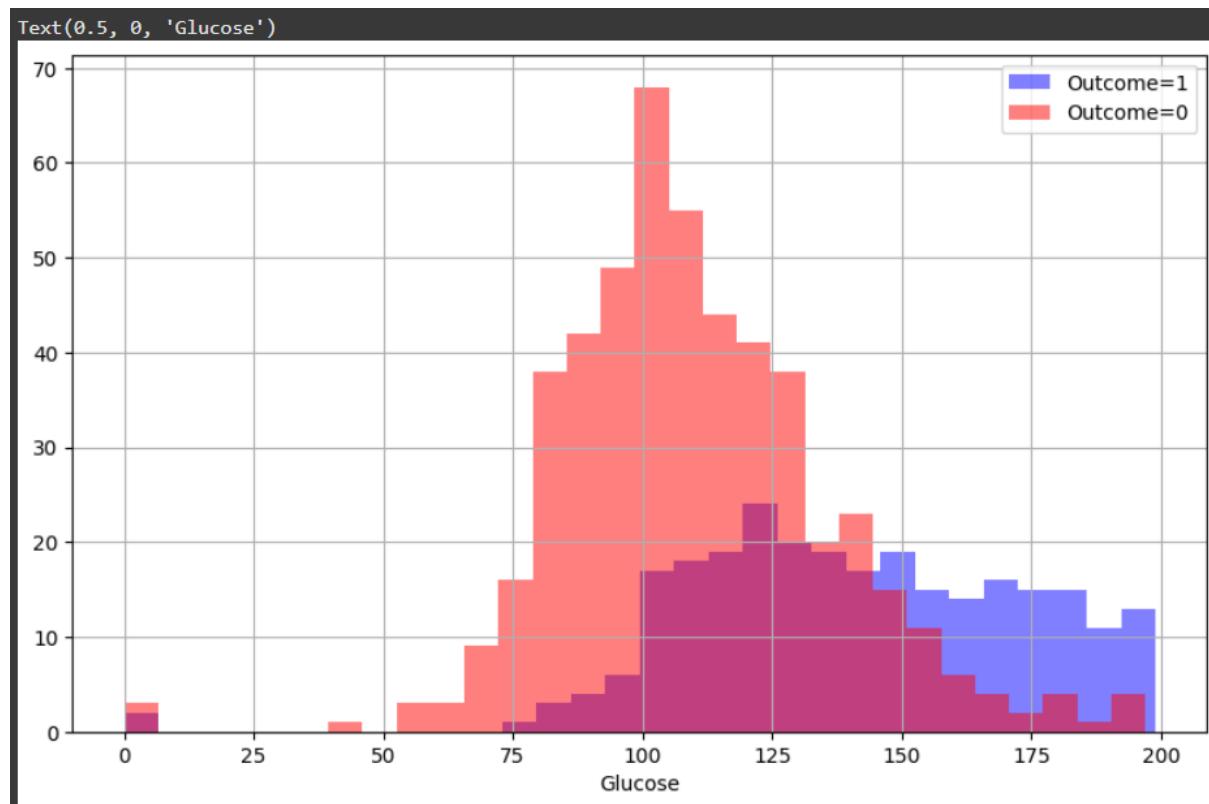

```

```
▶ diabetes.describe()  
→  


|       | Pregnancies | Glucose    | BloodPressure | SkinThickness | Insulin    | BMI        | DiabetesPedigreeFunction | Age        | Outcome    |
|-------|-------------|------------|---------------|---------------|------------|------------|--------------------------|------------|------------|
| count | 768.000000  | 768.000000 | 768.000000    | 768.000000    | 768.000000 | 768.000000 | 768.000000               | 768.000000 | 768.000000 |
| mean  | 3.845052    | 120.894531 | 69.105469     | 20.536458     | 79.799479  | 31.992578  | 0.471876                 | 33.240885  | 0.348958   |
| std   | 3.369578    | 31.972618  | 19.355807     | 15.952218     | 115.244002 | 7.884160   | 0.331329                 | 11.760232  | 0.476951   |
| min   | 0.000000    | 0.000000   | 0.000000      | 0.000000      | 0.000000   | 0.000000   | 0.078000                 | 21.000000  | 0.000000   |
| 25%   | 1.000000    | 99.000000  | 62.000000     | 0.000000      | 0.000000   | 27.300000  | 0.243750                 | 24.000000  | 0.000000   |
| 50%   | 3.000000    | 117.000000 | 72.000000     | 23.000000     | 30.500000  | 32.000000  | 0.372500                 | 29.000000  | 0.000000   |
| 75%   | 6.000000    | 140.250000 | 80.000000     | 32.000000     | 127.250000 | 36.600000  | 0.626250                 | 41.000000  | 1.000000   |
| max   | 17.000000   | 199.000000 | 122.000000    | 99.000000     | 846.000000 | 67.100000  | 2.420000                 | 81.000000  | 1.000000   |

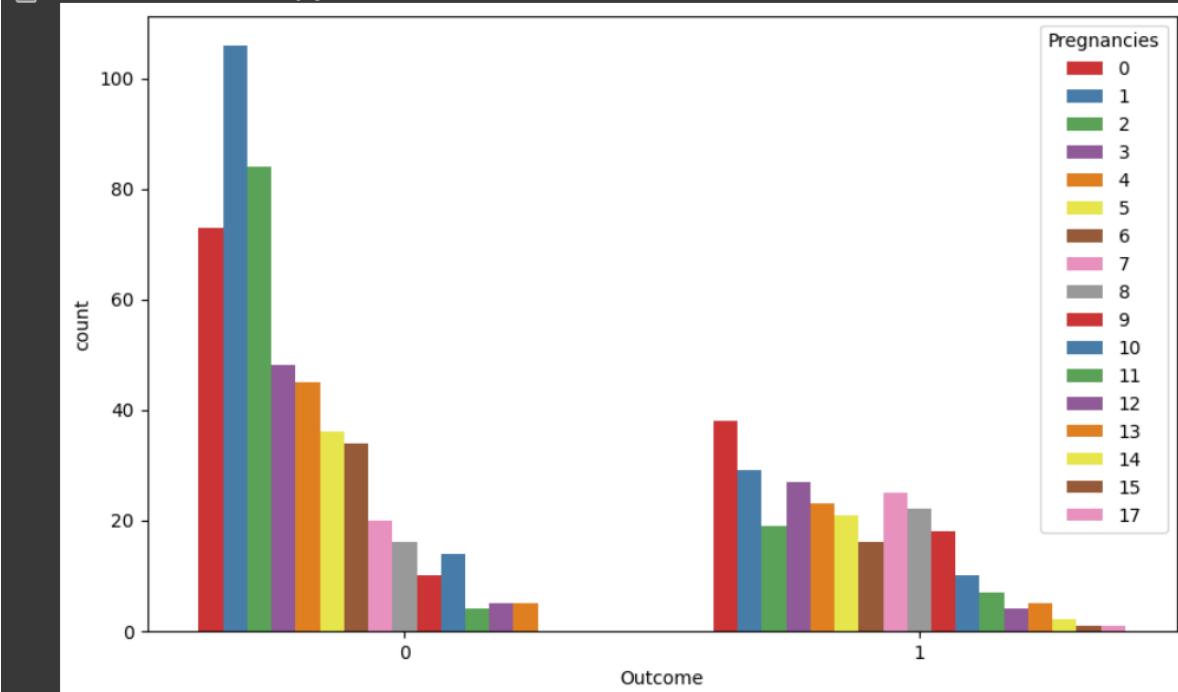

```

```
[ ] plt.figure(figsize=(10,6))  
diabetes[diabetes['Outcome']==1]['Glucose'].hist(alpha=0.5,color='blue',bins=30,label='Outcome=1')  
diabetes[diabetes['Outcome']==0]['Glucose'].hist(alpha=0.5,color='red',bins=30,label='Outcome=0')  
plt.legend()  
plt.xlabel('Glucose')
```



```
▶ plt.figure(figsize=(10,6))
sns.countplot(x='Outcome',hue='Pregnancies',data=diabetes,palette='Set1')
```

```
→ <Axes: xlabel='Outcome', ylabel='count'>
```



Saurav Malik

CSE-A

06017702722

**LEARNING OUTCOMES:**

---

---

---

---

## EXPERIMENT 2

**AIM:** Program to demonstrate simple Linear Regression

### **THEORY:**

Linear regression is one of the simplest and most widely used **statistical and machine learning techniques** for modeling the relationship between a **dependent variable (target)** and one or more **independent variables (predictors/features)**.

At its core, it tries to fit a **straight line (linear function)** to the data that best explains the relationship between the input(s) and the output.

### **Key Idea:**

- If you have **one independent variable**, the model is called **Simple Linear Regression**. The equation looks like:

$$y = \beta_0 + \beta_1 x + \epsilon$$

where:

- $y$  = dependent variable (output)
- $x$  = independent variable (input)
- $\beta_0$  = intercept (value of  $y$  when  $x=0$ )
- $\beta_1$  = slope (change in  $y$  for a unit change in  $x$ )
- $\epsilon$  = error term (captures randomness/unexplained part)

- If you have **multiple independent variables**, it's **Multiple Linear Regression**:

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n + \epsilon$$

### **Assumptions of Linear Regression:**

1. Linearity – Relationship between predictors and target is linear.
2. Independence – Observations are independent of each other.
3. Homoscedasticity – Constant variance of errors.
4. Normality – Errors are normally distributed (for inference).
1. No/low multicollinearity – Predictors should not be highly correlated

### **CODE:**

```
▶ import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
from sklearn.model_selection import train_test_split

▶ from collections.abc import ValuesView
ds=pd.read_csv('house.csv')
X=ds.iloc[:, :-1].values
Y=ds.iloc[:, 1].values

[ ] ds.head()

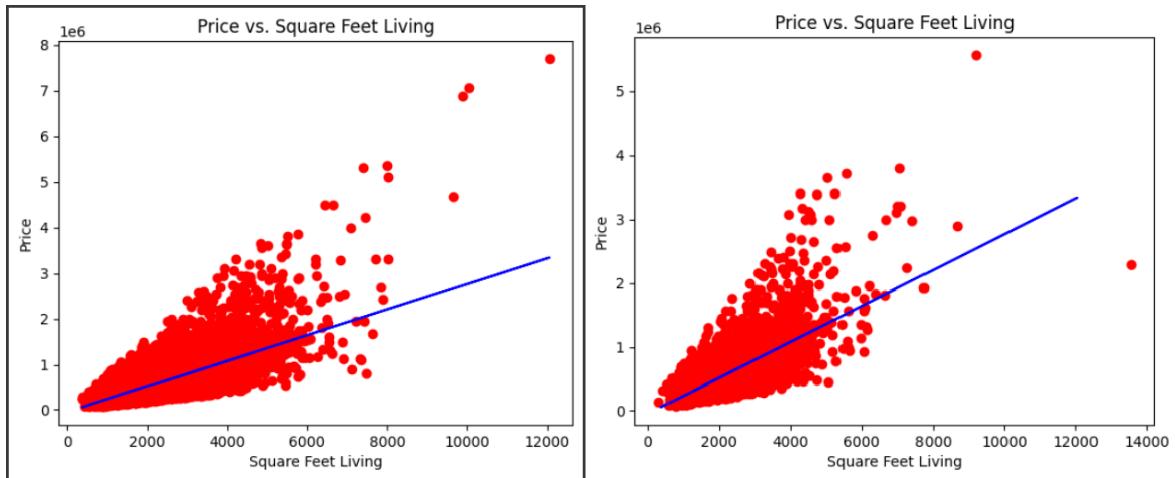
Σ sqft_living     price
0      1180  221900
1      2570  538000
2       770  180000
3      1960  604000
4      1680  510000
```

```
[ ] from sklearn.model_selection import train_test_split
X_train,X_test,Y_train,Y_test=train_test_split(X,Y,test_size=0.25)

from sklearn.linear_model import LinearRegression
regressor=LinearRegression()
regressor.fit(X_train,Y_train)

Σ ▾ LinearRegression ⓘ ⓘ
LinearRegression()

▶ Y_pred=regressor.predict(X_test)
plt.scatter(X_train,Y_train,color='red')
plt.plot(X_train,regressor.predict(X_train),color='blue')
plt.title('Price vs. Square Feet Living')
plt.xlabel('Square Feet Living')
plt.ylabel('Price')
plt.show()
plt.scatter(X_test,Y_test,color='red')
plt.plot(X_train,regressor.predict(X_train),color='blue')
plt.title('Price vs. Square Feet Living')
plt.xlabel('Square Feet Living')
plt.ylabel('Price')
plt.show()
```



## LEARNING OUTCOMES:

---

---

---

---

## EXPERIMENT 3

**AIM:** Program to demonstrate simple Logistic Regression

### THEORY:

Logistic Regression is a supervised machine learning algorithm mainly used for **classification problems**. Unlike linear regression, which predicts continuous values, logistic regression predicts the **probability of belonging to a class**. It first computes a linear combination of inputs:

$$z = w_0 + w_1 x_1 + w_2 x_2 + \dots + w_n x_n$$

and then applies the **sigmoid (logistic) function** to map the result between 0 and 1:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

Based on a threshold (commonly 0.5), the model decides the class, e.g., if  $\sigma(z) > 0.5$ , predict class 1; otherwise, class 0.

There are different types of logistic regression: **binary** (two classes), **multinomial** (more than two categories), and **ordinal** (ordered categories). Logistic regression is widely applied in domains like **medical diagnosis** (disease/no disease), **finance** (loan default prediction), **marketing** (customer churn), and **NLP** (sentiment analysis). In essence, it is a simple yet powerful algorithm for handling classification tasks by estimating probabilities and converting them into class labels.

## CODE:

```
[ ] import pandas as pd
col_names = ['pregnant', 'glucose', 'bp', 'skin', 'insulin', 'bmi', 'pedigree', 'age', 'label']
# load dataset
pima = pd.read_csv("pima-indians-diabetes.csv", header=None, names=col_names)

▶ pima.head()

→ pregnant glucose bp skin insulin bmi pedigree age label
0 6 148 72 35 0 33.6 0.627 50 1
1 1 85 66 29 0 26.6 0.351 31 0
2 8 183 64 0 0 23.3 0.672 32 1
3 1 89 66 23 94 28.1 0.167 21 0
4 0 137 40 35 168 43.1 2.288 33 1
```

+ Code + Text

```
[ ] #split dataset in features and target variable
feature_cols = ['pregnant', 'insulin', 'bmi', 'age','glucose','bp','pedigree']
X = pima[feature_cols] # Features
y = pima.label # Target variable

[ ] # split X and y into training and testing sets
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.35,random_state=0)

[ ] from sklearn.linear_model import LogisticRegression
logreg = LogisticRegression()
logreg.fit(X_train,y_train)
y_pred=logreg.predict(X_test)
```

/usr/local/lib/python3.12/dist-packages/sklearn/linear\_model/\_logistic.py:465: ConvergenceWarning:  
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

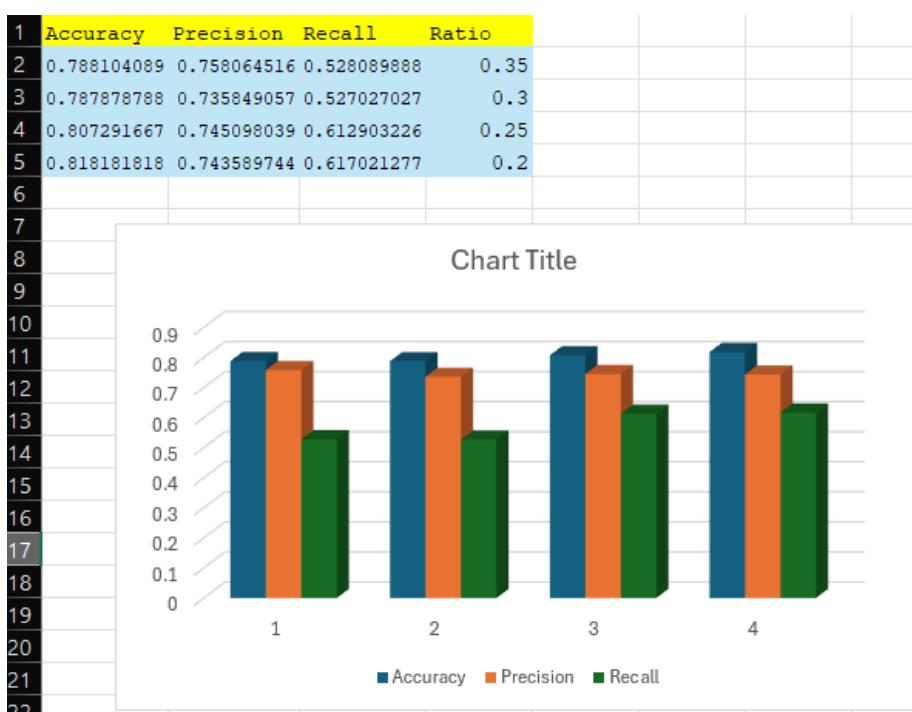
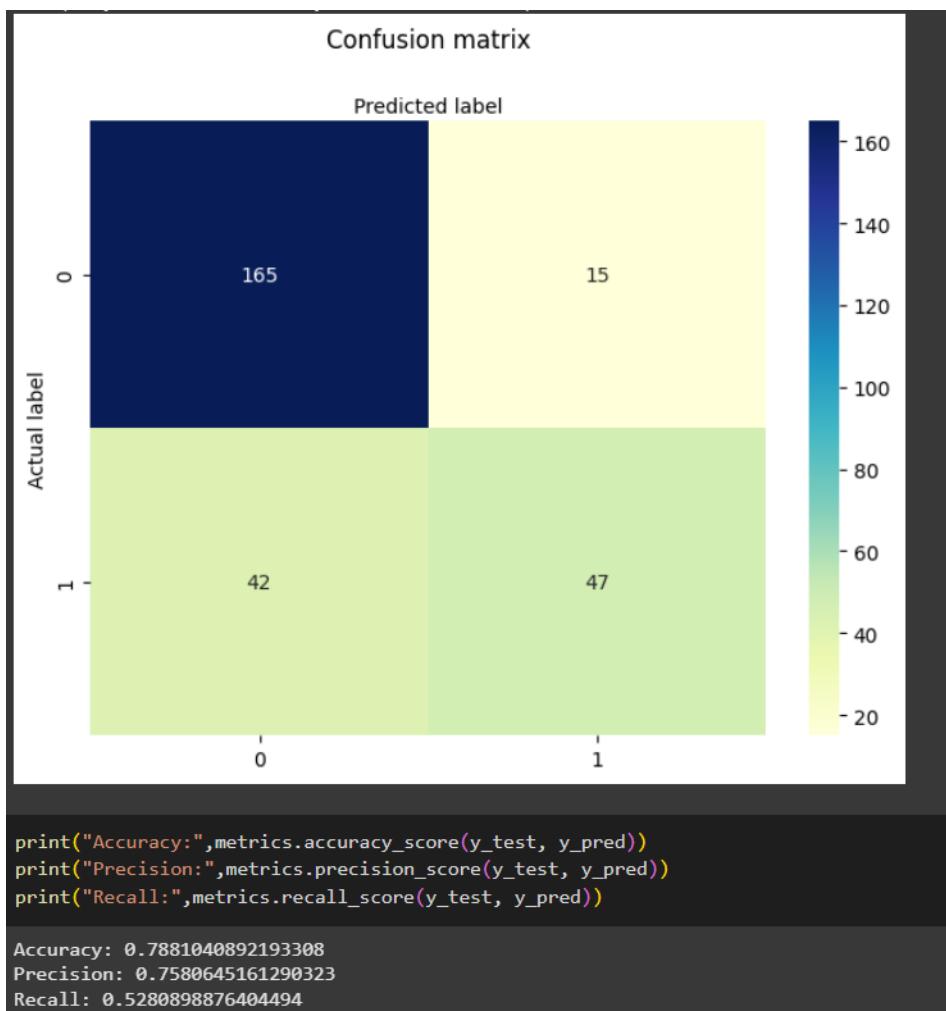
→ Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
 Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
`n_iter_i = _check_optimize_result()`

```
[ ] from sklearn import metrics
cnf_matrix = metrics.confusion_matrix(y_test, y_pred)
cnf_matrix
```

→ array([[165, 15],
 [ 42, 47]])

```
[ ] import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
▶ class_names=[0,1] # name of classes
fig, ax = plt.subplots()
tick_marks = np.arange(len(class_names))
plt.xticks(tick_marks, class_names)
plt.yticks(tick_marks, class_names)
# create heatmap
sns.heatmap(pd.DataFrame(cnf_matrix), annot=True, cmap="YlGnBu" ,fmt='g')
ax.xaxis.set_label_position("top")
plt.tight_layout()
plt.title('Confusion matrix', y=1.1)
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
```



Saurav Malik

CSE-A

06017702722

**LEARNING OUTCOMES:**

---

---

---

---

## EXPERIMENT 4

**Aim:** Program to demonstrate Decision Tree- ID3 Algorithm

**Theory:** A Decision Tree is a supervised learning method used for classification. It splits data into branches based on attributes until a decision (class label) is reached. The ID3 (Iterative Dichotomiser 3) algorithm, proposed by Ross Quinlan, builds trees using Entropy and Information Gain.

Key Concepts:

- Entropy (measure of impurity):

$$Entropy(S) = - \sum p_i \log_2(p_i)$$

where  $p_i$  is the probability of class  $i$

- Information Gain (attribute selection):

$$Gain(S, A) = Entropy(S) - \sum_v \frac{|S_v|}{|S|} \times Entropy(S_v)$$

Attribute with maximum gain is chosen for

splitting. Steps of ID3:

1. Calculate entropy of dataset.
2. Compute information gain for each attribute.
3. Select attribute with highest gain as decision node.
4. Split dataset into subsets and repeat recursively.
5. Stop when all records belong to one class or no attributes remain.

Advantages:

- Simple and easy to interpret.
- Works well with categorical data.

Limitations:

- Biased towards attributes with many values.
- Sensitive to noise and overfitting.

**CODE:**

```
import pandas
as pd import
numpy as np
from sklearn.preprocessing import LabelEncoder
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.metrics import confusion_matrix,
ConfusionMatrixDisplay import matplotlib.pyplot as plt

#
#
Da
ta
se
t
#
data = {
    'Outlook': ['Sunny', 'Sunny', 'Overcast', 'Rain', 'Rain', 'Rain',
    'Overcast',
        'Sunny', 'Sunny', 'Rain', 'Sunny', 'Overcast',
    'Overcast',
    'Rain'],
    'Temperature': ['Hot', 'Hot', 'Hot', 'Mild', 'Cool', 'Cool',
    'Cool',
        'Mild', 'Cool', 'Mild', 'Mild', 'Mild', 'Hot',
    'Mild'],
    'Humidity': ['High', 'High', 'High', 'High', 'Normal', 'Normal',
    'Normal',
        'High', 'Normal', 'Normal', 'Normal', 'High',
    'Normal', 'High'],
    'Wind': ['Weak', 'Strong', 'Weak', 'Weak', 'Weak', 'Strong',
    'Strong',
        'Weak', 'Weak', 'Weak', 'Strong', 'Strong', 'Weak',
    'Strong'],
    'PlayTennis': ['No', 'No', 'Yes', 'Yes', 'Yes', 'No', 'Yes',
    'No', 'Yes', 'Yes', 'Yes', 'Yes', 'Yes', 'Yes', 'No']
}

df = pd.DataFrame(data)

#
# Label
Encoding
#
encoder =
LabelEncoder()
for col in
df.columns:
```

```
df[col] = encoder.fit_transform(df[col])

X =
df.drop('PlayTennis',
axis=1) y =
df['PlayTennis']

#
#
Decision
Tree #
#
clf = DecisionTreeClassifier(criterion='entropy',
random_state=0) clf.fit(X, y)

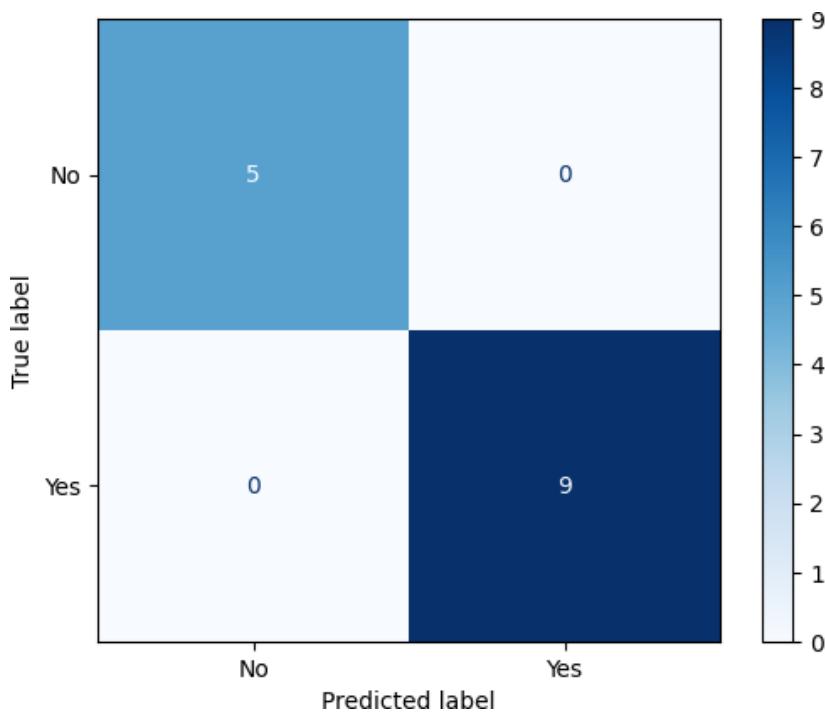
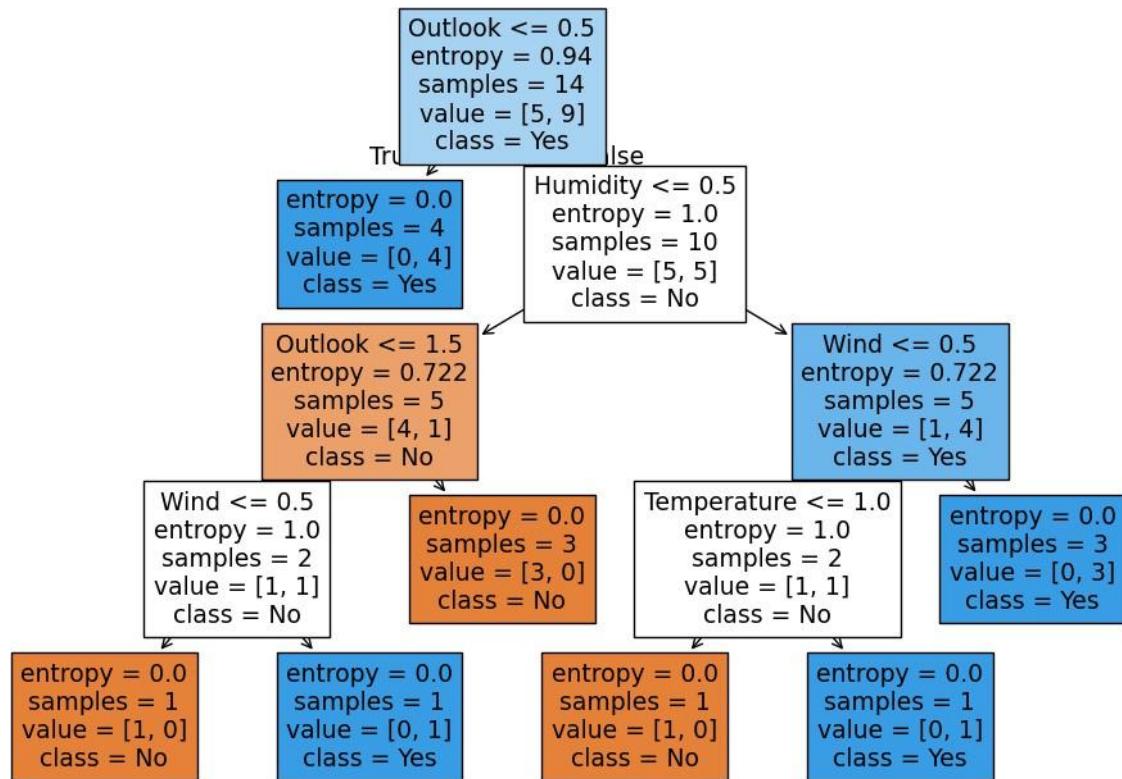
#
# Plot
Decision
Tree #
plt.figure(figsize=(12,8))
plot_tree(clf, feature_names=X.columns, class_names=['No', 'Yes'],
filled=True) plt.show()

#
#
Prediction
Example #
sample = [[0, 1, 0, 1]]
print("Predicted:",
clf.predict(sample))

#
#
Confusion
Matrix #
y_pred = clf.predict(X)
cm = confusion_matrix(y, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm,
display_labels=['No', 'Yes']) disp.plot(cmap="Blues")
plt.show()

print("\n--- Confusion Matrix Values
---") print("Confusion Matrix:\n",
cm)
tn, fp, fn, tp = cm.ravel()
print(f"True Negative: {tn}, False Positive: {fp}, False Negative:
{fn}, True Positive: {tp}")
```

```
#  
# Entropy &  
Information Gain #  
def entropy(labels):  
    """Calculate entropy of a list of labels"""  
    values, counts = np.unique(labels,  
    return_counts=True) probs = counts /  
    counts.sum()  
    return -np.sum(probs * np.log2(probs))  
  
def information_gain(data, feature, target):  
    """Calculate Information Gain of a feature relative to the  
target"""  
    total_entropy = entropy(data[target])  
    values, counts = np.unique(data[feature], return_counts=True)  
    weighted_entropy = 0  
    for v, c in zip(values, counts):  
        subset = data[data[feature] == v][target]  
        weighted_entropy += (c/len(data)) *  
        entropy(subset)  
    return total_entropy - weighted_entropy  
  
# Dataset before encoding for readability  
df_original = pd.DataFrame(data)  
  
print("\n--- Entropy & Information Gain ---")  
print(f"Entropy of PlayTennis:  
{entropy(df_original['PlayTennis']):.4f}")  
  
for feature in ['Outlook', 'Temperature',  
    'Humidity', 'Wind']: ig =  
    information_gain(df_original, feature, 'PlayTennis')
```

**OUTPUT:**

```
--- Confusion Matrix Values ---
Confusion Matrix:
[[5 0]
 [0 9]]
True Negative: 5, False Positive: 0, False Negative: 0, True Positive: 9

--- Entropy & Information Gain ---
Entropy of PlayTennis: 0.9403
Information Gain(Outlook): 0.2467
Information Gain(Temperature): 0.0292
Information Gain(Humidity): 0.1518
Information Gain(Wind): 0.0481
```

## Learning Outcomes:

---

---

---

---